Monitoring

English version of ../../operations/monitoring.md.

Signal	Where	Purpose
/health	api	liveness, used by the image HEALTHCHECK
/metrics	api, requires AEZA_METRICS_TOKEN	latency and codes per route (backend/app/metrics.py)
Structured logs	container stdout	incident analysis
Snapshot provenance	domain/provenance.py	understand which data produced a number
Model state	/api/v1/model/*	availability of the LLM provider

Minimal alerts

1. /health does not answer for 2 minutes.
2. The 5xx share is above 1% over 5 minutes.
3. /metrics is unreachable or the token is not set in production.
4. The snapshot has not been refreshed for longer than expected (AEZA_SNAPSHOT_POLL_SECONDS).
5. LLM provider errors with AEZA_AI_ENABLED=1 — feature degradation, not a service outage.

What is missing

- No distributed tracing: one service, but background recalculations (AEZA_RECALCULATE_SECONDS) are not linked to requests.
- No SLO expressed as a number.
- No out-of-the-box alerts — metrics are exposed, the rules belong to operations.